

IDEB

4 de Setembro de 2017

1 Example: IDEB Analysis

We are going to analyse data corresponding to the IDEB (Basic Education Development Index) for brazilian cities. The data comes from the file

```
In [1]: ideb_file = "IDEB por Município Rede Federal Séries Finais (51 a 81).xml"
```

which was obtained from the main brazilian government open data site dados.gov.br

Since we have an .xml file, we'll use the *xml.etree.ElementTree* module to parse its contents. For simplicity, we'll call this module *ET*.

```
In [2]: import xml.etree.ElementTree as ET
```

1.0.1 The ElementTree (ET) module

An XML file is a hierarchical set of data, so the most intuitive way to represent this data is by a tree. To do this, the ET module implements two classes: the *ElementTree* class represents the whole XML file as a tree, and the *Element* class represents one node of this tree. All interactions that occur with the whole file (like reading and writing to this file) are done through the *ElementTree* class; on the other hand, every interaction with an isolated element of the XML and its subelements are done through the *Element* class.

By reading the docs, we learn that the *ET.parse* methods returns an *ElementTree* from a file.

```
In [3]: tree = ET.parse(ideb_file)
```

The *ElementTree* class has the following structure:

```
In [4]: dir(tree)
```

```
Out[4]: ['__class__',
         '__delattr__',
         '__dict__',
         '__dir__',
         '__doc__',
         '__eq__',
         '__format__',
         '__ge__',
         '__getattribute__',
         '__gt__',
```

```

'__hash__',
'__init__',
'__init_subclass__',
'__le__',
'__lt__',
'__module__',
'__ne__',
'__new__',
'__reduce__',
'__reduce_ex__',
'__repr__',
'__setattr__',
'__sizeof__',
'__str__',
'__subclasshook__',
'__weakref__',
'_root',
'_setroot',
'find',
'findall',
'findtext',
'getiterator',
'getroot',
'iter',
'iterfind',
'parse',
'write',
'write_c14n']

```

According to the documentation for this module, we access the ElementTree via its *root* node, which is an *Element* class instance. To see the root element, we use the *getroot* method:

```
In [5]: root = tree.getroot()
```

As an *Element*, the root object has the *tag* and *attrib* properties, and *attrib* is a dictionary of its attributes. Let's see what are these values:

```
In [6]: root.tag
```

```
Out[6]: 'result'
```

```
In [7]: root.attrib
```

```
Out[7]: {}
```

To access each child node of the root element, we iterate on these nodes (which are also *Elements*):

```
In [8]: for child in root:
        print(child.tag, child.attrib)
```

```
url {}
id {}
nome {}
nome_estendido {}
descricao {}
inicio {}
final {}
formatacao {}
data_atualizacao {}
aditividade {}
url_origem {}
tempo_aditividade {}
portal_dados_abertos {}
disponibilizacao {}
estado {}
fonte_gestora {}
fonte_provedora {}
grupo_informacao {}
base_territorial {}
periodicidade {}
multiplicador {}
produto {}
publicacao {}
unidade_medida {}
orgao_primeiro_escalao {}
valores {}
```

We can see that our XML comes with a lot of data. Next, we will try to get a subset of this data.

1.0.2 Selecting the data

Now that we have a better idea of the document's structure, let's build a pandas *DataFrame* with what we need. First, we can see that we only need the last node of the root element, "valores"(which stands for "values" in Portuguese); the other nodes are in fact just the header for the XML file. Let's explore this node.

```
In [9]: IDEBvalues = root.find('valores')
```

Note that there is one more layer of data here:

```
In [10]: IDEBvalues
```

```
Out[10]: <Element 'valores' at 0x7f159c12f598>
```

```
In [11]: IDEBvalues[0]
```

```
Out[11]: <Element 'entry' at 0x7f159c12f5e8>
```

Now, we can explore the grandchildren of the root node:

[illegible]

[illegible]

```

ano {}
valor {}
municipio_ibge {}
ano {}
valor {}
municipio_ibge {}
ano {}
valor {}
municipio_ibge {}
ano {}
valor {}
municipio_ibge {}
ano {}
valor {}
municipio_ibge {}
ano {}
valor {}
municipio_ibge {}
ano {}
valor {}
municipio_ibge {}
ano {}

```

Now, let's extract the data we are interested in:

```

In [13]: data = []
         for child in IDEBvalues:
             data.append([float(child[0].text), child[1].text, child[2].text])

```

```

In [14]: data

```

```

Out[14]: [[4.7, '120040', '2009'],
          [6.0, '130260', '2009'],
          [5.9, '140010', '2009'],
          [5.5, '150140', '2009'],
          [4.0, '211130', '2009'],
          [6.3, '220190', '2009'],
          [6.9, '230440', '2009'],
          [7.1, '261160', '2009'],
          [6.5, '280670', '2009'],
          [7.1, '292740', '2009'],
          [6.0, '310620', '2009'],
          [6.3, '313670', '2009'],
          [6.4, '317020', '2009'],
          [5.7, '330455', '2009'],

```

```
[6.9, '410690', '2009'],
[5.7, '420540', '2009'],
[5.8, '431490', '2009'],
[7.3, '431690', '2009'],
[7.1, '500270', '2009'],
[5.3, '520870', '2009'],
[4.4, '120040', '2007'],
[5.3, '140010', '2007'],
[5.2, '150140', '2007'],
[3.4, '211130', '2007'],
[4.3, '220190', '2007'],
[6.8, '230440', '2007'],
[7.5, '261160', '2007'],
[5.4, '280670', '2007'],
[7.2, '292740', '2007'],
[5.5, '310620', '2007'],
[7.0, '313670', '2007'],
[6.0, '317020', '2007'],
[6.1, '330455', '2007'],
[5.8, '420540', '2007'],
[6.2, '431490', '2007'],
[6.5, '431690', '2007'],
[6.5, '500270', '2007'],
[5.5, '520870', '2007'],
[6.7, '530010', '2007']]
```

Since Pandas seems to be fashionable right now ;) let's use it to store and treat this data. We'll give it a shorter name though, pd.

```
In [15]: import pandas as pd
```

Now, we create our DataFrame from the preexisting data.

```
In [16]: IDEBTable = pd.DataFrame(data, columns = ["Valor", "Municipio", "Ano"])
```

```
In [17]: IDEBTable
```

```
Out[17]:
```

	Valor	Municipio	Ano
0	4.7	120040	2009
1	6.0	130260	2009
2	5.9	140010	2009
3	5.5	150140	2009
4	4.0	211130	2009
5	6.3	220190	2009
6	6.9	230440	2009
7	7.1	261160	2009
8	6.5	280670	2009
9	7.1	292740	2009
10	6.0	310620	2009

11	6.3	313670	2009
12	6.4	317020	2009
13	5.7	330455	2009
14	6.9	410690	2009
15	5.7	420540	2009
16	5.8	431490	2009
17	7.3	431690	2009
18	7.1	500270	2009
19	5.3	520870	2009
20	4.4	120040	2007
21	5.3	140010	2007
22	5.2	150140	2007
23	3.4	211130	2007
24	4.3	220190	2007
25	6.8	230440	2007
26	7.5	261160	2007
27	5.4	280670	2007
28	7.2	292740	2007
29	5.5	310620	2007
30	7.0	313670	2007
31	6.0	317020	2007
32	6.1	330455	2007
33	5.8	420540	2007
34	6.2	431490	2007
35	6.5	431690	2007
36	6.5	500270	2007
37	5.5	520870	2007
38	6.7	530010	2007

You can see there are two sets of data here, one for 2007 and another for 2009. We'll only use the most recent data for our "analysis".

```
In [18]: IDEBTable = IDEBTable.loc[0:19]
```

1.0.3 Identifying the city codes

In our IDEBTable, cities are identified by their so called "IBGE Code", which is a code issued to each locality by the Brazilian Institute for Geography and Statistics (IBGE). In order to make this more user friendly, we'll read the most recent Excel file with the list of cities and their respective 7 digit codes (from 2014; these codes include a final verification digit). For this, we'll use the xlrld module, which must be manually installed; see this.

```
In [19]: localCodesIBGE = pd.read_excel("DTB_2014_Municipio.xls")
```

Now we can inspect the data by using the pandas *head* method for DataFrames:

```
In [20]: localCodesIBGE.head()
```



```

Out[20]:
  UF  Nome_UF  Mesorregião Geográfica  Nome_Mesorregião \
0  11  Rondônia                    2  Leste Rondoniense
1  11  Rondônia                    2  Leste Rondoniense
2  11  Rondônia                    2  Leste Rondoniense
3  11  Rondônia                    2  Leste Rondoniense
4  11  Rondônia                    2  Leste Rondoniense

  Microrregião Geográfica  Nome_Microrregião  Município \
0                    6          Cacoal          15
1                    3          Ariquemes         23
2                    8  Colorado do Oeste         31
3                    6          Cacoal          49
4                    8  Colorado do Oeste         56

  Cod Municipio Completo  Nome_Município
0                    1100015  Alta Floresta D'Oeste
1                    1100023          Ariquemes
2                    1100031          Cabixi
3                    1100049          Cacoal
4                    1100056        Cerejeiras

```

The columns we are interested in are just "Nome_UF", "Cod Municipio Completo" and "Nome_Município", which stand for State (or Province), Complete City Code and City Name, respectively.

```
In [21]: localCodesIBGE = localCodesIBGE[["Nome_UF", "Cod Municipio Completo", "Nome_Município"]]
```

Now, we have two DataFrames: **IDEBTable**, containing the complete IDEB data corresponding to city names, and **localCodesIBGE**, containing the corresponding city codes. We must select from the complete **localCodesIBGE** table only the rows corresponding to cities for which we have the IDEB value. For this, we will extract from both DataFrames the columns corresponding to the city codes (remember that in the **localCodesIBGE** table, codes have an extra verification code which we will not use):

```
In [42]: IDEBCities = IDEBTable["Município"]
        cities = localCodesIBGE["Cod Municipio Completo"].map(lambda x: str(x)[0:6])
```

Note that we have used *map* to transform numerical data into strings, removing the last digit.

Now, both **IDEBCities** and **cities** are pandas Series objects. To get the indices of cities for which we have IDEB data, first we will identify which codes are **not** in **IDEBCities**:

```
In [43]: citiesToRemove = cities[~cities.isin(IDEBCities)]
```

We remove the corresponding rows from the **localCodesIBGE** table:

```
In [45]: newTable = localCodesIBGE.drop(citiesToRemove.index).reset_index(drop=True)
```

Finally, we will create a new DataFrame joining city name and IDEB value:

```
In [46]: finalData = pd.concat([newTable, IDEBTable], axis=1)
```

This gives

In [47]: finalData

```
Out[47]:
```

	Nome_UF	Cod Municipio	Completo	Nome_Município	Valor	\
0	Acre		1200401	Rio Branco	4.7	
1	Amazonas		1302603	Manaus	6.0	
2	Roraima		1400100	Boa Vista	5.9	
3	Pará		1501402	Belém	5.5	
4	Maranhão		2111300	São Luís	4.0	
5	Piauí		2201903	Bom Jesus	6.3	
6	Ceará		2304400	Fortaleza	6.9	
7	Pernambuco		2611606	Recife	7.1	
8	Sergipe		2806701	São Cristóvão	6.5	
9	Bahia		2927408	Salvador	7.1	
10	Minas Gerais		3106200	Belo Horizonte	6.0	
11	Minas Gerais		3136702	Juiz de Fora	6.3	
12	Minas Gerais		3170206	Uberlândia	6.4	
13	Rio de Janeiro		3304557	Rio de Janeiro	5.7	
14	Paraná		4106902	Curitiba	6.9	
15	Santa Catarina		4205407	Florianópolis	5.7	
16	Rio Grande do Sul		4314902	Porto Alegre	5.8	
17	Rio Grande do Sul		4316907	Santa Maria	7.3	
18	Mato Grosso do Sul		5002704	Campo Grande	7.1	
19	Goiás		5208707	Goiânia	5.3	

	Municipio	Ano
0	120040	2009
1	130260	2009
2	140010	2009
3	150140	2009
4	211130	2009
5	220190	2009
6	230440	2009
7	261160	2009
8	280670	2009
9	292740	2009
10	310620	2009
11	313670	2009
12	317020	2009
13	330455	2009
14	410690	2009
15	420540	2009
16	431490	2009
17	431690	2009
18	500270	2009
19	520870	2009

1.1 Finishing up: a pretty figure

In order to include graphics in notebooks, usually the first cell in the notebook contains the code

```
% matplotlib inline  
or  
% matplotlib notebook
```

Since we don't want to sacrifice the legibility of our *article* by starting it with some mysterious command, we can use the `init_cell` nbextension so that a later cell is executed first on our notebook (Section ??).

First, let's import the pyplot sublibrary of the matplotlib library and call it `plt`:

```
In [48]: import matplotlib.pyplot as plt
```

We'll do a very simple plot, but for this it would be nice to use the city names instead of the numerical indices in the `finalData` table:

```
In [49]: finalData.set_index(["Nome_Município"], inplace=True)
```

Now, we will select the column with the values ("Valor") for the IDEB by city in the `finalData` table (note that the result of this operation is a Series):

```
In [50]: finalData["Valor"]
```

```
Out [50]: Nome_Município  
Rio Branco      4.7  
Manaus          6.0  
Boa Vista       5.9  
Belém           5.5  
São Luís        4.0  
Bom Jesus       6.3  
Fortaleza       6.9  
Recife          7.1  
São Cristóvão   6.5  
Salvador        7.1  
Belo Horizonte  6.0  
Juiz de Fora    6.3  
Uberlândia      6.4  
Rio de Janeiro  5.7  
Curitiba       6.9  
Florianópolis   5.7  
Porto Alegre    5.8  
Santa Maria     7.3  
Campo Grande    7.1  
Goiânia         5.3  
Name: Valor, dtype: float64
```

We are ready for our pretty (yet irrelevant) picture.

```
In [52]: finalData["Valor"].plot(kind='barh')  
plt.title("IDEB by city (Data from 2009)")
```

```
Out [52]: <matplotlib.text.Text at 0x7f156c4e7e48>
```

1.2 Comments about automatic documentation and script generation

To convert this notebook to a regular Python .py script, use

```
jupyter-nbconvert --to python 'IDEB.ipynb' --template=removeextracode.tpl
```

The removeextracode.tpl has the following content:

```
{% extends 'python.tpl'%}

{% block input %}
{% if 'codecomment' in cell['metadata'].get('tags', []) %}
    {{ cell.source | comment_lines }}
{% else %}
    {{ cell.source | ipython2python }}
{% endif %}
{% endblock input %}
```

This means that we will include all notebook cells tagged with **codecomment** as comments on our script. This is to avoid generating a unusable script including our inspection of objects and attempts at solving a problem.

For more details on templates and the nbconvert extension, check this page, for example.

1.3 Initialization cell

Through the "init_cell" extension (also from nbextensions), it is possible to alter the order of execution of notebook cells. If we look at the metadata of the cell below, we can see that it is marked to be executed before all other cells, and so we obtain the desired result when we run all cells in the notebook. (This command allows us to see inline graphics inside our notebook).

```
In [53]: %matplotlib inline
```